

# AI Medical Chatbot Pro — LLM + LangChain + Pinecone + Flask/Streamlit

---

An end-to-end **AI-powered Medical Chatbot** built using **Large Language Models (Llama 3.3 70B)**, **LangChain**, **Pinecone vector search**, and **Flask/Streamlit** — deployable on both **AWS (EC2 + ECR)** and **Render.com**.

This project covers every part:

- Data ingestion & processing
  - Embedding + Vector DB with Pinecone
  - RAG-based chatbot pipeline using LangChain
  - REST API with Flask + Beautiful UI
  - Streamlit version for easy deployment
  - Docker containerization
  - Multiple deployment options (AWS EC2, Render)
  - Automated CI/CD with GitHub Actions
- 

## Features

-  **Powered by Llama 3.3 70B** via Groq (Free API)
-  **Advanced RAG System** - Retrieves relevant medical information
-  **Streaming Responses** - Real-time answer generation
-  **Two UI Options** - Flask (advanced) or Streamlit (simple)
-  **Responsive Design** - Works on all devices
-  **Privacy First** - No conversation storage
-  **Free to Deploy** - Multiple free hosting options

## Live Demo

- **Render Deployment:** [Your Render URL]
- **AWS Deployment:** [Your AWS URL]

## Screenshots

[Deployment Page](#)

## CI/CD Pipeline

## Quick Start - Local Development

### STEP 01 — Clone the Repository

```
git clone https://github.com/SunnyPanchani/AI-Medical-Chatbot-Pro.git
cd AI-Medical-Chatbot-Pro
```

## STEP 02 — Create Virtual Environment

### Using Conda:

```
conda create -n medibot python=3.11 -y  
conda activate medibot
```

### Or using venv:

```
python -m venv env  
source env/bin/activate # On Windows: env\Scripts\activate
```

## STEP 03 — Install Dependencies

```
pip install -r requirements.txt
```

## STEP 04 — Set Up Environment Variables

Create a `.env` file in the root directory:

```
# Pinecone Vector Database  
PINECONE_API_KEY=xxxxxxxxxxxxxxxxxxxxxxxxxxxxxx  
  
# Groq LLM API (FREE - get from https://console.groq.com)  
GROQ_API_KEY=gsk_xxxxxxxxxxxxxxxxxxxxxxxxxxxxxx
```

### Get Your Free API Keys:

- **Pinecone:** <https://www.pinecone.io/> (Free tier: 100K vectors)
- **Groq:** <https://console.groq.com/> (Free tier: Very generous limits)

## STEP 05 — Create Vector Store

**Important:** Run this once to create your Pinecone index:

```
python store_index.py
```

This will:

- Load PDFs from `data/` folder
- Create embeddings
- Store vectors in Pinecone

- Create index: `medical-chatbot`

## STEP 06 — Test the Setup

```
python test_chatbot.py
```

This runs test queries to verify everything works.

## STEP 07 — Run the Application

### Option A: Flask Application (Advanced UI)

```
python app.py
```

Then open: `http://localhost:8080`

### Option B: Streamlit Application (Simple UI)

```
streamlit run streamlit_app.py
```

Then open: `http://localhost:8501`

---

## 🔧 Tech Stack

Component	Technology	Purpose
<b>LLM</b>	Llama 3.3 70B (Groq)	Answer generation
<b>Framework</b>	LangChain	RAG pipeline
<b>Vector DB</b>	Pinecone	Semantic search
<b>Embeddings</b>	Sentence Transformers	Text embeddings
<b>Backend</b>	Flask	REST API
<b>UI</b>	Flask + Streamlit	User interfaces
<b>Deployment</b>	AWS EC2, Render	Cloud hosting
<b>CI/CD</b>	GitHub Actions	Automation
<b>Containerization</b>	Docker	Packaging

---

## 📁 Project Structure

```

AI Medical Chatbot Pro/
├── .github/
│   └── workflows/
│       └── cicd.yaml      # CI/CD pipeline
├── .streamlit/
│   └── config.toml      # Streamlit config
└── data/
    └── Medical_book.pdf  # Medical knowledge base
src/
├── __init__.py
├── helper.py           # Core functions
└── prompt.py           # Prompt templates
static/
├── style.css            # Flask UI styles
└── script.js            # Flask UI scripts
templates/
└── index.html          # Flask UI template
research/
└── trails.ipynb        # Development notebook
app.py                  # Flask application
streamlit_app.py        # Streamlit application
store_index.py          # Vector store creation
test_chatbot.py         # Test script
requirements.txt        # Dependencies
render.yaml             # Render config
Dockerfile              # Docker config
setup.py                # Package setup
.env                    # Environment variables (not in git)
.gitignore              # Git ignore
README.md               # This file

```

## Deployment Options

Option 1: Deploy on Render (Easiest - Free Tier Available)

**Perfect for:** Quick deployment, testing, demos

### 1. Prerequisites:

- GitHub account
- Render account (free)
- API keys (Pinecone, Groq)

### 2. Deploy Steps:

```

# 1. Create vector store locally (one time only)
python store_index.py

# 2. Push to GitHub

```

```
git add .
git commit -m "Deploy to Render"
git push

# 3. On Render.com:
- New Web Service
- Connect GitHub repository
- Add environment variables:
  * PINECONE_API_KEY
  * GROQ_API_KEY
- Deploy!
```

**Your app will be live at:** <https://your-app.onrender.com>

**Detailed Guide:** See [RENDER\\_DEPLOYMENT.md](#)

### Render Features:

- Free tier available
- Auto-deploy on git push
- HTTPS included
- Easy setup (5 minutes)

Option 2: Deploy on AWS EC2 (Production - Full Control)

**Perfect for:** Production, custom domains, full control

### Prerequisites

#### 1. AWS Account

#### 2. IAM User with Permissions:

- [AmazonEC2ContainerRegistryFullAccess](#)
- [AmazonEC2FullAccess](#)

### Deployment Steps

#### 1. Login to AWS Console

#### 2. Create IAM User for Deployment

Create user with these policies:

- [AmazonEC2ContainerRegistryFullAccess](#)
- [AmazonEC2FullAccess](#)

#### 3. Create ECR Repository

```
# Create ECR repository to store Docker image
Repository Name: medicalbot
Region: us-east-1
```

```
# Save the URI:  
994626601219.dkr.ecr.us-east-1.amazonaws.com/medicalbot
```

## 4. Create EC2 Instance

- AMI: Ubuntu 22.04 LTS
- Instance Type: t2.medium (recommended) or t2.micro (free tier)
- Storage: 20 GB
- Security Group: Allow ports 22 (SSH), 80 (HTTP), 8080 (App)

## 5. Install Docker on EC2

SSH into your EC2 instance:

```
# Update system  
sudo apt-get update -y  
sudo apt-get upgrade -y  
  
# Install Docker  
curl -fsSL https://get.docker.com -o get-docker.sh  
sudo sh get-docker.sh  
  
# Add user to docker group  
sudo usermod -aG docker ubuntu  
newgrp docker  
  
# Verify Docker installation  
docker --version
```

## 6. Configure EC2 as Self-Hosted Runner

In your GitHub repository:

1. Go to: Settings → Actions → Runners
2. Click: New self-hosted runner
3. Choose: Linux
4. Run commands on EC2 instance

## 7. Setup GitHub Secrets

Add these secrets in: Repository Settings → Secrets and variables → Actions

```
AWS_ACCESS_KEY_ID=your_access_key  
AWS_SECRET_ACCESS_KEY=your_secret_key  
AWS_DEFAULT_REGION=us-east-1  
ECR_REPO=994626601219.dkr.ecr.us-east-1.amazonaws.com/medicalbot  
PINECONE_API_KEY=your_pinecone_key  
GROQ_API_KEY=your_groq_key
```

## 8. Create Vector Store

**Important:** Before deploying, create the vector store:

```
# Run locally or on EC2
python store_index.py
```

## 9. Deploy with GitHub Actions

```
# Push to main branch
git add .
git commit -m "Deploy to AWS"
git push origin main

# GitHub Actions will automatically:
# 1. Build Docker image
# 2. Push to ECR
# 3. Deploy to EC2
```

## 10. Access Your Application

```
http://your-ec2-public-ip:8080
```

## Docker Commands

### Build Docker Image

```
docker build -t medicalbot .
```

### Run Docker Container

```
docker run -p 8080:8080 \
-e PINECONE_API_KEY=your_key \
-e GROQ_API_KEY=your_key \
medicalbot
```

### Push to ECR

```
# Login to ECR
aws ecr get-login-password --region us-east-1 | \
```

```
docker login --username AWS --password-stdin \
994626601219.dkr.ecr.us-east-1.amazonaws.com

# Tag image
docker tag medicalbot:latest \
994626601219.dkr.ecr.us-east-1.amazonaws.com/medicalbot:latest

# Push image
docker push 994626601219.dkr.ecr.us-east-1.amazonaws.com/medicalbot:latest
```

---

## 🧪 Testing

### Run Test Script

```
python test_chatbot.py
```

### Test in Notebook

```
jupyter notebook research/trails.ipynb
```

### Manual Testing

```
from src.helper import setup_medical_chatbot, ask_question

# Initialize
rag_chain, retriever, docsearch = setup_medical_chatbot()

# Ask question
result = ask_question(
    rag_chain,
    "What is diabetes?",
    retriever=retriever
)
```

---

## 🔧 Configuration

### Adjust Model Parameters

Edit [src/helper.py](#):

```
chatModel = initialize_groq_llm(
    model_name="llama-3.3-70b-versatile", # Model choice
```

```
    temperature=0.3,  
    max_tokens=1024  
)  
  
# 0 = focused, 1 = creative  
# Response length
```

## Change Retrieved Documents

```
rag_chain, retriever = create_rag_chain(  
    docsearch,  
    chatModel,  
    num_documents=5 # Number of sources (3-7 recommended)  
)
```

## Customize Prompts

Edit [src/prompt.py](#) to modify system prompts.

## Performance Metrics

Metric	Value
Average Response Time	2-5 seconds
Vectors in Database	5,859 chunks
Documents Retrieved	5 per query
Model	Llama 3.3 70B
Embedding Dimension	384
Cost	\$0 (Free APIs)

## Security & Privacy

-  No conversation storage
-  API keys in environment variables
-  HTTPS supported (on Render/AWS)
-  No user tracking
-  Data processed in memory only
-  .env file excluded from git

## Important Notes

### Medical Disclaimer

This chatbot provides **informational content only**. It is **NOT a substitute for professional medical advice, diagnosis, or treatment**. Always seek the advice of your physician or other qualified health provider with any

questions you may have regarding a medical condition.

## API Usage

- **Groq:** Free tier with generous limits
- **Pinecone:** Free tier (100K vectors, 1 index)
- Monitor your usage in respective dashboards

## Data Privacy

- Conversations are NOT stored
- Data is processed in real-time
- No personal information is collected

---

## 🐛 Troubleshooting

### Common Issues

#### 1. ModuleNotFoundError

```
pip install -r requirements.txt
```

#### 2. Pinecone Index Not Found

```
python store_index.py
```

#### 3. GROQ\_API\_KEY Not Found

- Check `.env` file exists
- Verify key is correct
- Restart application

#### 4. Out of Memory (Render Free Tier)

- Upgrade to paid tier
- Or reduce `num_documents` to 3

#### 5. Slow Response

- Check internet connection
- Verify API keys are valid
- Try reducing retrieved documents

---

## 🤝 Contributing

Contributions are welcome! Please:

1. Fork the repository
  2. Create a feature branch (`git checkout -b feature/YourFeature`)
  3. Commit changes (`git commit -m 'Add YourFeature'`)
  4. Push to branch (`git push origin feature/YourFeature`)
  5. Open a Pull Request
- 

## License

This project is licensed under the MIT License - see [LICENSE](#) file.

---

## Acknowledgments

- **Groq** - For free Llama 3.3 70B API access
  - **Pinecone** - For vector database
  - **LangChain** - For RAG framework
  - **Streamlit** - For UI framework
  - **Flask** - For web framework
- 

## Contact

### Sunny Panchani

- GitHub: [@SunnyPanchani](#)
  - Email: [sunnypanchani9007@gmail.com](mailto:sunnypanchani9007@gmail.com)
  - LinkedIn: [Your LinkedIn]
- 

## Star This Project

If you find this project helpful, please give it a star! 

---

## Roadmap

- Basic chatbot with RAG
  - Flask web interface
  - Streamlit interface
  - Docker deployment
  - AWS deployment
  - Render deployment
  - Multi-language support
  - Voice input/output
  - Mobile app version
  - Chat history export
  - More medical specialties
-

## Tips for Best Results

1. **Ask Specific Questions:** "What are the symptoms of Type 2 diabetes?" instead of just "diabetes"
  2. **Check Sources:** Review cited documents for verification
  3. **Consult Professionals:** Always verify with healthcare providers
  4. **Report Issues:** Use GitHub issues for bugs or suggestions
- 

## Documentation

- [Render Deployment Guide](#)
  - [Deployment Checklist](#)
  - [Streaming Guide](#)
  - [API Documentation](#) (coming soon)
- 

Made with ❤️ by Sunny Panchani

★ If this project helped you, please star it! ★